

---

## 3. Time and Ordering

Sistemes Distribuïts en Xarxa (SDX)  
Facultat d'Informàtica de Barcelona (FIB)  
Universitat Politècnica de Catalunya (UPC)  
2024/2025 Q2

# Contents

---

- **Introduction**

- Physical clocks

- Logical clocks

# Time

---

- Time is unambiguous in centralized systems
  - System clock keeps time, all entities use this
- No global time on distributed systems
  - Each node has a (crystal-based) system clock
    - Less accurate than atomic clocks
    - Results in **clock skew** (two clocks, two times) and **clock drift** (two clocks, two count rates)
    - Drifts 1 second every 11 days w.r.t. a perfect clock
  - Problem: An event that occurred after another may be assigned an earlier time
    - Use physical and logical clocks to deal with this

# Contents

---

- Introduction

- **Physical clocks**

- Logical clocks

# Physical clocks


---

- Physical clocks allow to synchronize nodes ...
  - i. with a master node (with a UTC receiver)
    - UTC: Universal Coordinated Time is an international standard based on atomic time that is broadcasted through short-wave radio and satellite
  - ii. with one another
- ... within a given bound
  - **Perfect** clock synchronization is not feasible
    - Synchronization limited by network jitter and clock drift
    - Typical accuracy of milliseconds
  - Clocks approximate real time but cannot generally be used to find out the **order** of events

# Physical clocks


- Synchronize at least every  $R < \delta/2\rho$  to limit skew between two clocks to less than  $\delta$  time units

$R$ : resynchronization interval  
 $\rho$ : maximum clock drift rate  
 $\delta$ : maximum allowed clock skew



# Cristian's algorithm

- Synchronize nodes with a server with UTC receiver within a given bound: **External synchronization**
  - Intended for intranets with a UTC-sync server
- 1. Each client asks the time to the server at every  $R$  interval
- 2. Client sets the time to  $T_s + \text{RTT}/2$ 
  - RTT: round-trip time
    - $T_2 - T_1$
  - Assumes symmetrical latency
  - Accuracy of client's clock is  $\pm(\text{RTT}/2 - \text{Min})$
- NTP is based on a similar concept

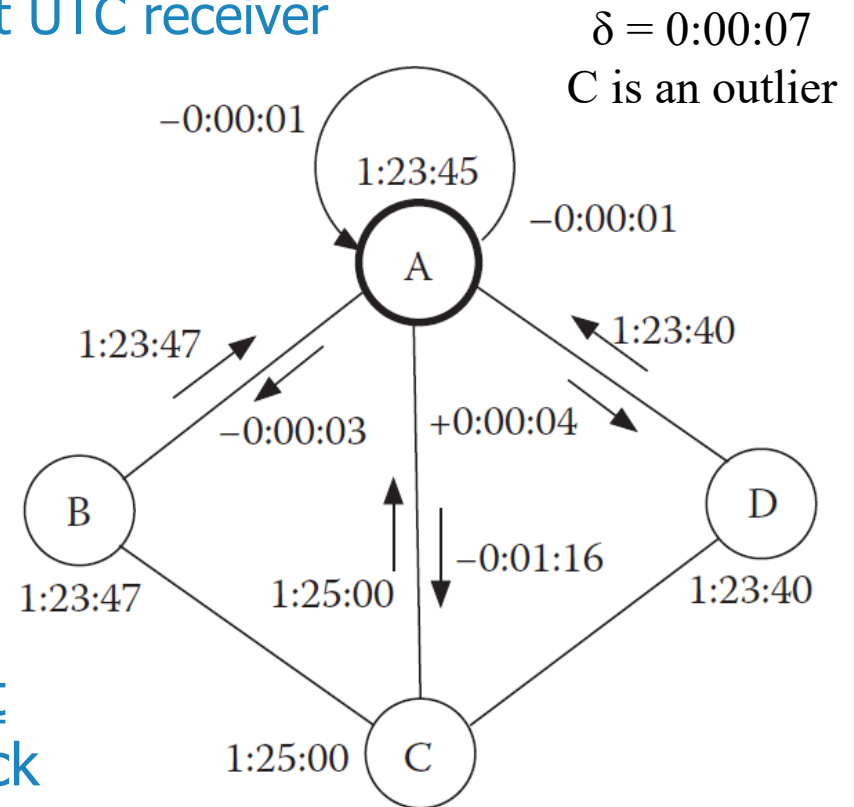


# Berkeley algorithm

- Keep clocks synchronized with one another within a given bound: **Internal synchronization**

- Intended for intranets without UTC receiver

- Master polls the clocks of all the slaves at every  $R$  interval
  - Adapts them by considering round-trip times
- Master calculates a fault-tolerant average
  - Clocks lying outside the given bound are discarded
- Master sends the adjustment to be made to each local clock





# System clock discipline

---

- Time stepping can have side effects
  - i. When setting the time forward
    - Some time instants are lost
    - Events scheduled at these times can be affected
  - ii. When setting the time backward
    - Monotonicity (time always moves forward) is violated
    - Subsequent events can appear earlier than previous ones
- Workaround: Slew (change gradually) the time by speeding it up or slowing it down until the adjustment has been achieved

# READING REPORT

---

**[Neville-Neil16]** Neville-Neil, G.V., *Time Is an Illusion, Lunchtime Doubly So*, Communications of the ACM, Vol. 59, No. 1, pp. 50-55, January 2016

# Contents

---

- Introduction
- Physical clocks
- **Logical clocks**


# Happened-before relation

---

- Processes need to know if event 'a' happened before or after event 'b'
  - Agree on the **order** in which events occur rather than the **time** at which they occurred
- The happened-before relation
  - If **a** and **b** are two events in the same process, and **a** comes before **b**, then **a**  $\rightarrow$  **b**
  - If **a** is the sending of a message, and **b** is the receipt of that message, then **a**  $\rightarrow$  **b**
  - If **a**  $\rightarrow$  **b** and **b**  $\rightarrow$  **c**, then **a**  $\rightarrow$  **c**

# Happened-before relation

- Example



$a \rightarrow b, b \rightarrow c, c \rightarrow d, d \rightarrow f, a \rightarrow f$  but  $a || e$  (concurrent)

➤ Happened-before relation defines a **partial ordering**

# Logical clocks

---

- To capture the happened-before relation, we attach a timestamp  $C(e)$  to each event  $e$ , satisfying the following properties:
  1. If  $a$  and  $b$  are two events in the same process, and  $a \rightarrow b$ , then  $C(a) < C(b)$
  2. If  $a$  corresponds to sending a message  $m$ , and  $b$  to the receipt of  $m$ , then  $C(a) < C(b)$
- How to attach a timestamp to an event when there is no global clock?  
 $\Rightarrow$  Use Lamport's logical clocks


# Lamport's logical clocks

---

- Each process  $P_i$  maintains a local counter  $C_i$
- $C_i$  is used to attach a timestamp to each event at  $P_i$
- $C_i$  is adjusted according to the following rules:
  1. When an event happens at  $P_i$ , it increases  $C_i$  by 1
  2. When  $P_i$  sends message  $m$  to  $P_j$ , sets  $ts(m) = C_i$
  3. When  $P_j$  receives  $m$ , sets  $C_j = \max(C_j, ts(m))$ , and then increases by 1

# Lamport's logical clocks

- Example



Note that  $C(e) < C(b)$  but  $b \parallel e$




# Lamport's logical clocks

---

- Lamport's clocks define a partial order that is consistent with the happened-before relation
  - i.e. it is consistent with causal order
- What if we need totally-ordered clocks?
  - Use Lamport's clocks, but in case two of them are equal, process IDs will be used to break the tie
    - $\{C_i(a), i\} < \{C_j(b), j\}$  if:
      - $C_i(a) < C_j(b)$  OR
      - $C_i(a) = C_j(b)$  AND  $i < j$
  - Can be used for instance to order the entry of processes to a critical section

# Logical clocks

- Lamport's clocks don't guarantee that if  $C(a) < C(b)$  then 'a' causally preceded 'b' ( $a \rightarrow b$ )



- $C(a) < C(c)$ , and 'a' causally preceded 'c' ( $a \rightarrow c$ )
- $C(e) < C(c)$ , but 'e' did not causally precede 'c' ( $c \parallel e$ )

⇒ Use vector clocks

# Vector clocks

---

- Each process  $P_i$  has an array  $VC_i[1...N]$ 
  - $VC_i[j]$  denotes the number of events that process  $P_i$  knows have taken place at process  $P_j$ 
    - i.e.  $VC_i[j]$  is the logical clock of  $P_j$  at process  $P_i$
- VC is adjusted as follows:
  1. When  $P_i$  sends a message  $m$ , it adds 1 to  $VC_i[i]$ , and sends  $VC_i$  with  $m$  as vector timestamp  $ts(m)$
  2. When  $P_j$  receives a message  $m$  from  $P_i$ , it updates each  $VC_j[k]$  to  $\max(VC_j[k], ts(m)[k])$ , and then increments  $VC_j[j]$  by 1


# Vector clocks

---

- Vector clocks allow to detect causality
  - $\mathbf{VC(a)} < \mathbf{VC(b)} \Leftrightarrow \mathbf{a \rightarrow b}$ 
    - If neither  $\mathbf{VC(a)} < \mathbf{VC(b)}$  nor  $\mathbf{VC(b)} < \mathbf{VC(a)}$ , then  $\mathbf{a || b}$
- How to compare vector clocks?
  - $\mathbf{VC(a)} < \mathbf{VC(b)}$  if and only if:
    - $\forall k: 1, \dots, N: \mathbf{VC(a)[k]} \leq \mathbf{VC(b)[k]}$
    - AND
    - $\exists k: 1, \dots, N: \mathbf{VC(a)[k]} < \mathbf{VC(b)[k]}$

# Vector clocks

- Example



Neither  $VC(b) < VC(e)$  nor  $VC(e) < VC(b)$ , so  $b || e$

# Summary

---

- We can synchronize physical clocks, but only within a given bound
- We cannot in general use physical time to find out the order of events
- Use logical clocks to find out events ordering
  - Lamport's clocks:  $a \rightarrow b \Rightarrow C(a) < C(b)$
  - Vector clocks:  $a \rightarrow b \Leftrightarrow VC(a) < VC(b)$
- Further details:
  - [Tanenbaum]: chapters 5.1 and 5.2
  - [Coulouris]: chapter 14