

---

# 7. Distributed File Systems

Sistemes Distribuïts en Xarxa (SDX)  
Facultat d'Informàtica de Barcelona (FIB)  
Universitat Politècnica de Catalunya (UPC)  
2024/2025 Q2

# Contents

---

- **Architecture**
- Communication
- Naming
- Synchronization
- Consistency & replication
- Fault tolerance

# Contents

---

- **Architecture**
  - **Client-server architectures**
  - Cluster-based architectures
- Communication
- Naming
- Synchronization
- Consistency & replication
- Fault tolerance

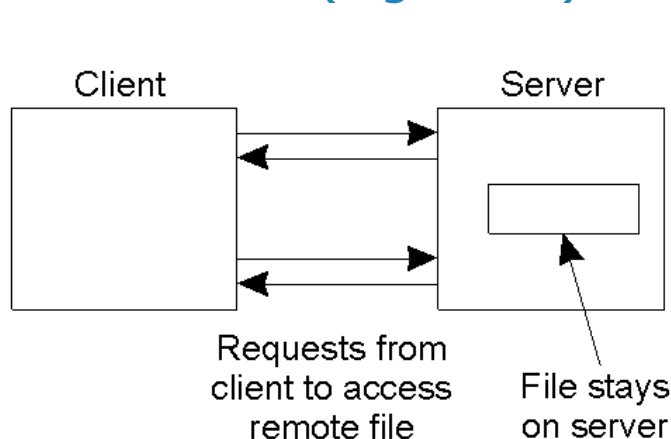
# Client-server architectures

## a) Remote access model (a.k.a. remote file service)


- Client is offered transparent access to a file system that is managed by a remote server (e.g. NFS)

## b) Upload/download model

- Client accesses locally to a file after downloading it from the server (e.g. Coda)



(a)



(b)

# Client-server architectures

---

## Remote access model

- Most of the work done at the server
- ↑ Consistency is easy
- ↑ Client does not need storage
- ↓ Server is a bottleneck
- ↓ Low tolerance to server crash and partitions
- Low communication cost to open files but high to operate on them

## Upload/download model

- Most of the work done at the client
- ↓ Consistency is hard
- ↓ Client downloads file: needs storage
- ↑ Load is distributed
- ↑ Tolerance to server crash and partitions
- High communication cost to open files but low to operate on them

# Client-server architectures

---

## Stateful server

- Server provides open and close operations
- Server keeps client state between requests on the same file

↑ Better performance

↓ Server state must be recovered after a crash

## Stateless server

- Server does not keep any client state between requests
  - A request must provide all the needed data to execute the operation

↓ Worse performance

↑ No need to restore any state after a crash

↓ Cannot be always followed to full extent

- e.g. file locking

# Sun Network File System (NFS)

---

- Designed for use on UNIX-based systems, but has been implemented for many OS
- Client-server (remote) file system
  - Each server provides a standardized view of its local file system, no matter how it is implemented
  - Clients access files via a communication protocol
- File system model similar to UNIX
  - Files treated as uninterpreted sequences of bytes
  - Files have name, but are referred by a file handle
- NFS was initially stateless, v4 is stateful


# NFS operations

Operation	V3	v4	Description
create	Yes	No	Create a regular file
create	No	Yes	Create a nonregular file (symbolic link, directory, special file)
link	Yes	Yes	Create a hard link to a file
symlink	Yes	No	Create a symbolic link to a file
mkdir	Yes	No	Create a subdirectory in a given directory
mknod	Yes	No	Create a special file
rename	Yes	Yes	Change the name of a file
remove	Yes	Yes	Remove a file from a file system
rmdir	Yes	No	Remove an empty subdirectory from a directory
open	No	Yes	Open a file (also create a regular file)
close	No	Yes	Close a file
lookup	Yes	Yes	Look up a file by means of a file name
readdir	Yes	Yes	Read the entries in a directory
readlink	Yes	Yes	Read the path name stored in a symbolic link
getattr	Yes	Yes	Read the attribute values for a file
setattr	Yes	Yes	Set one or more attribute values for a file
read	Yes	Yes	Read the data contained in a file
write	Yes	Yes	Write data to a file



# NFS architecture

- **Virtual File System (VFS)** provides uniform & transparent access to local and remote files



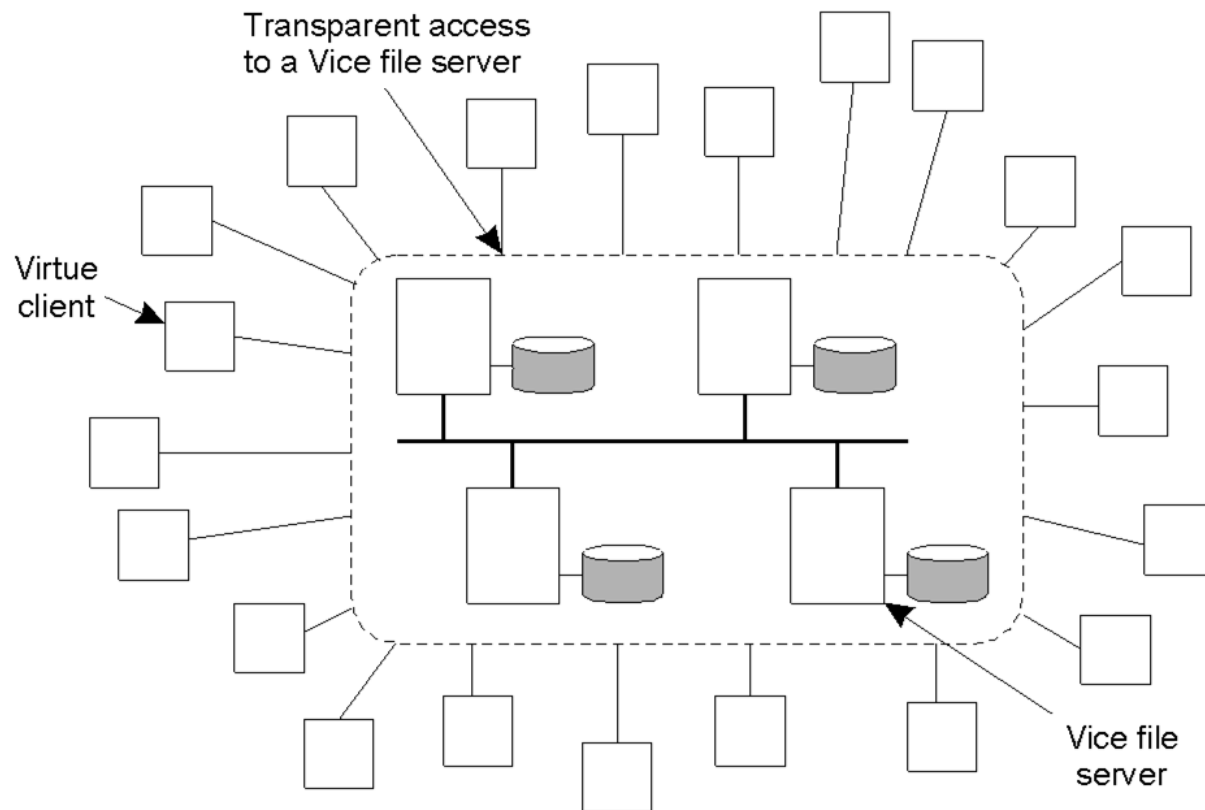
# Coda File System

---

- Based on Andrew File System (AFS v2)
  - Design philosophy: “Scalability and availability are more important than consistency”
    - Scalability: support a large number of users
    - Availability: tolerate server/communication failures and voluntary client disconnections (mobile users)
- ⇒ Use upload/download model
- Whole-file serving and caching
    - AFS v3 allows file data to be transferred and cached in 64-kbyte blocks

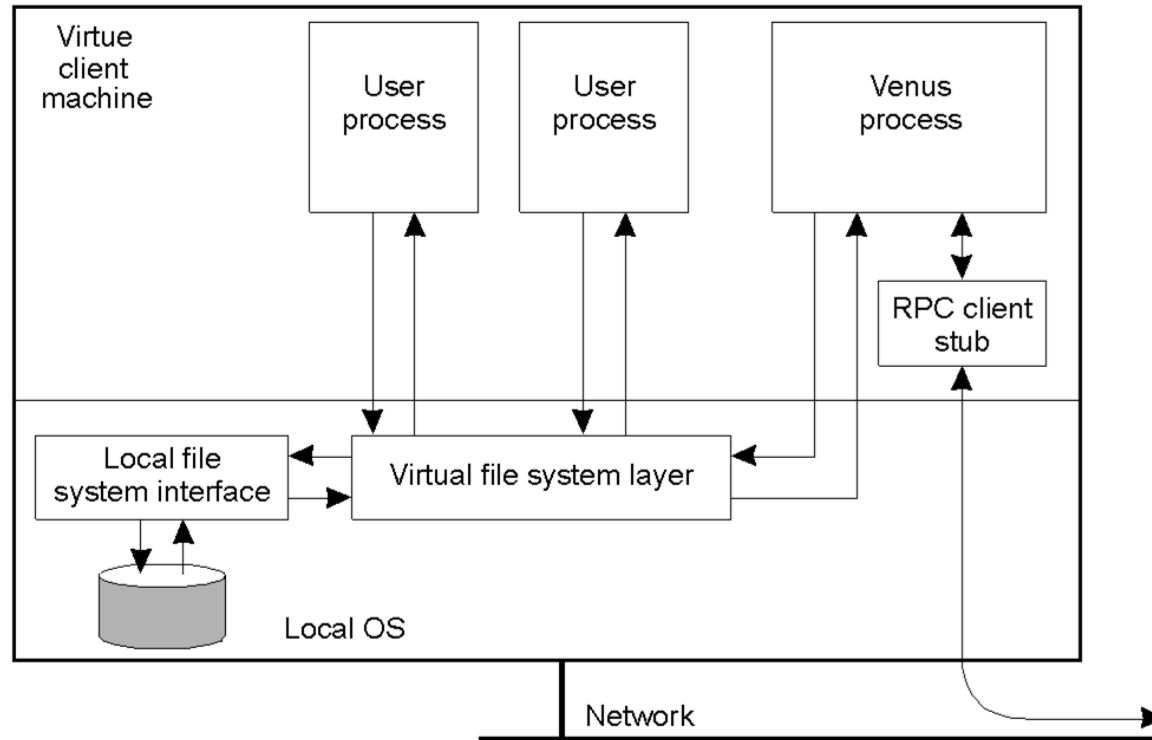
# Coda architecture

- Few centrally administered **Vice** file servers
- A large number of **Virtue** clients



# Coda architecture

- Clients host a user-level process **Venus** (~NFS client)
- Clients use VFS (as NFS does)
  - Coda looks like a traditional UNIX file system




# Contents

---

- **Architecture**
  - Client-server architectures
  - **Cluster-based architectures**
- Communication
- Naming
- Synchronization
- Consistency & replication
- Fault tolerance

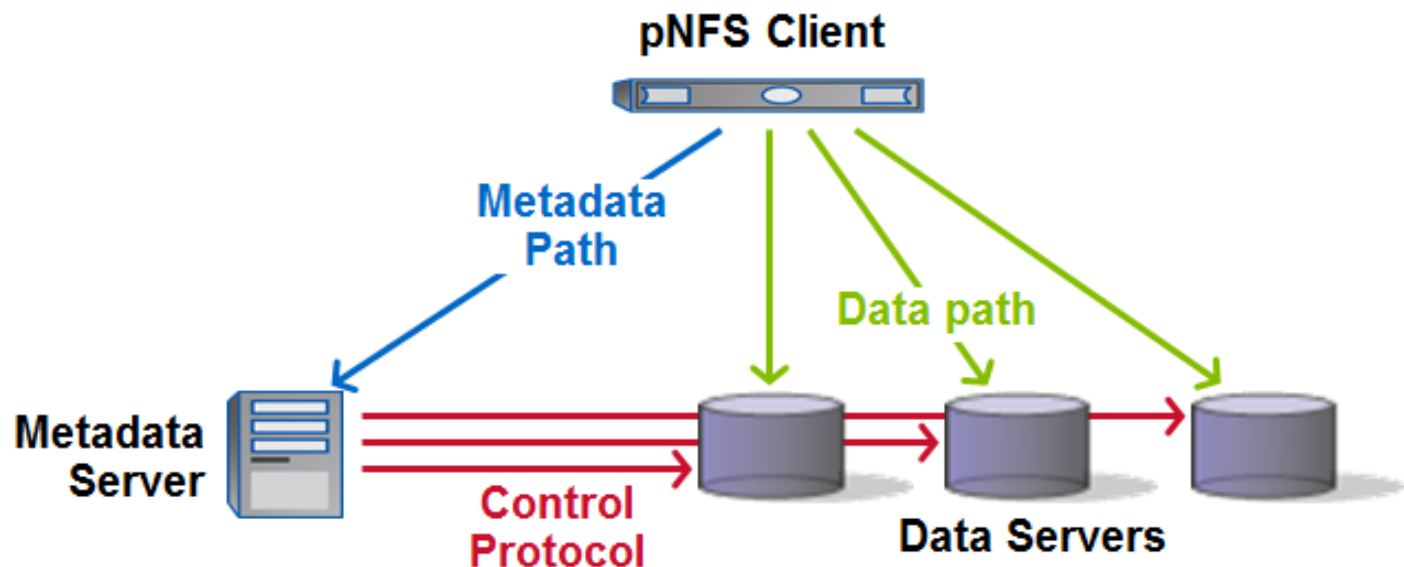
# Cluster-based architectures

- How to improve performance in server clusters with very large data collections?
- **File-stripping** techniques by which files can be fetched in parallel
  - A single file is distributed across multiple servers



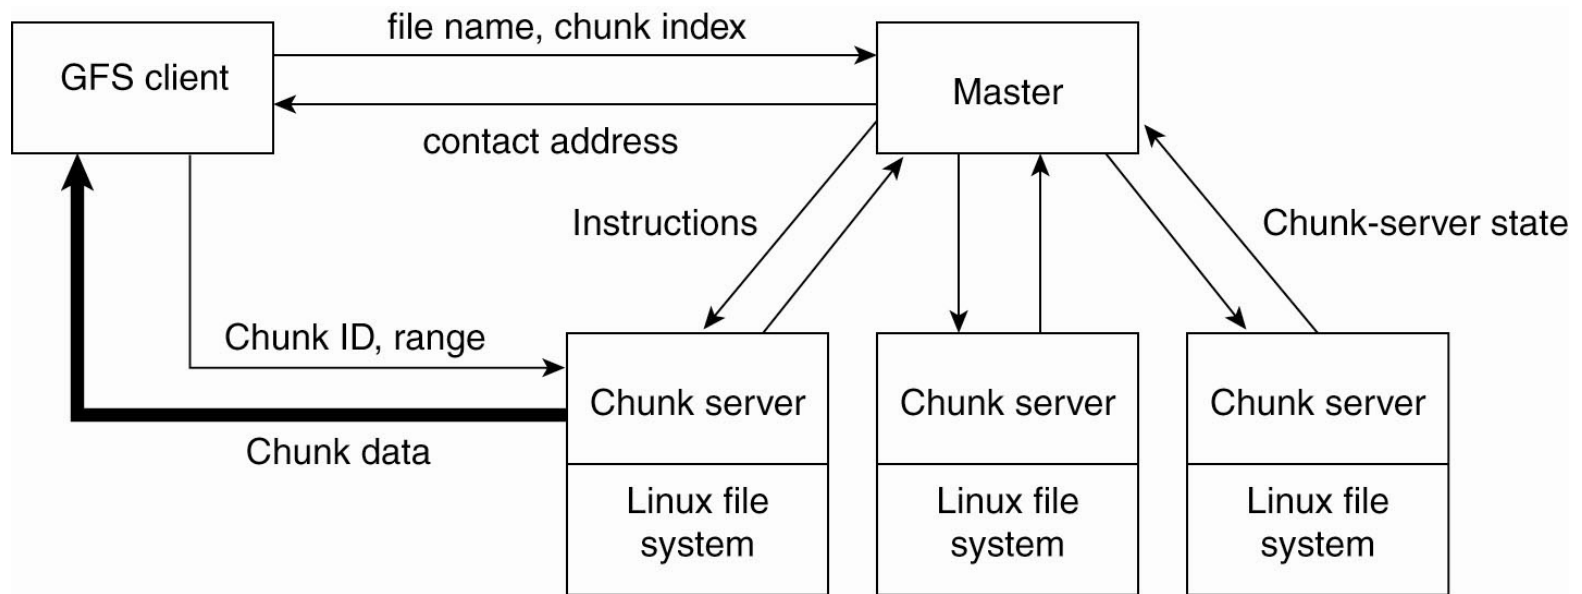
# Cluster-based architectures

- Ex: pNFS (part of NFS v4.1)
  - Metadata and data separation
    - Metadata server is out of the data path
  - Files are striped across a number of data servers



# Cluster-based architectures

- Ex: Google File System (GFS)
  - Each GFS cluster: 1 master, N chunk servers
  - Files divided into 64 MB chunks, which are replicated and distributed across chunk servers






# Contents

---

- Architecture
- **Communication**
- Naming
- Synchronization
- Consistency & replication
- Fault tolerance

# NFS communication

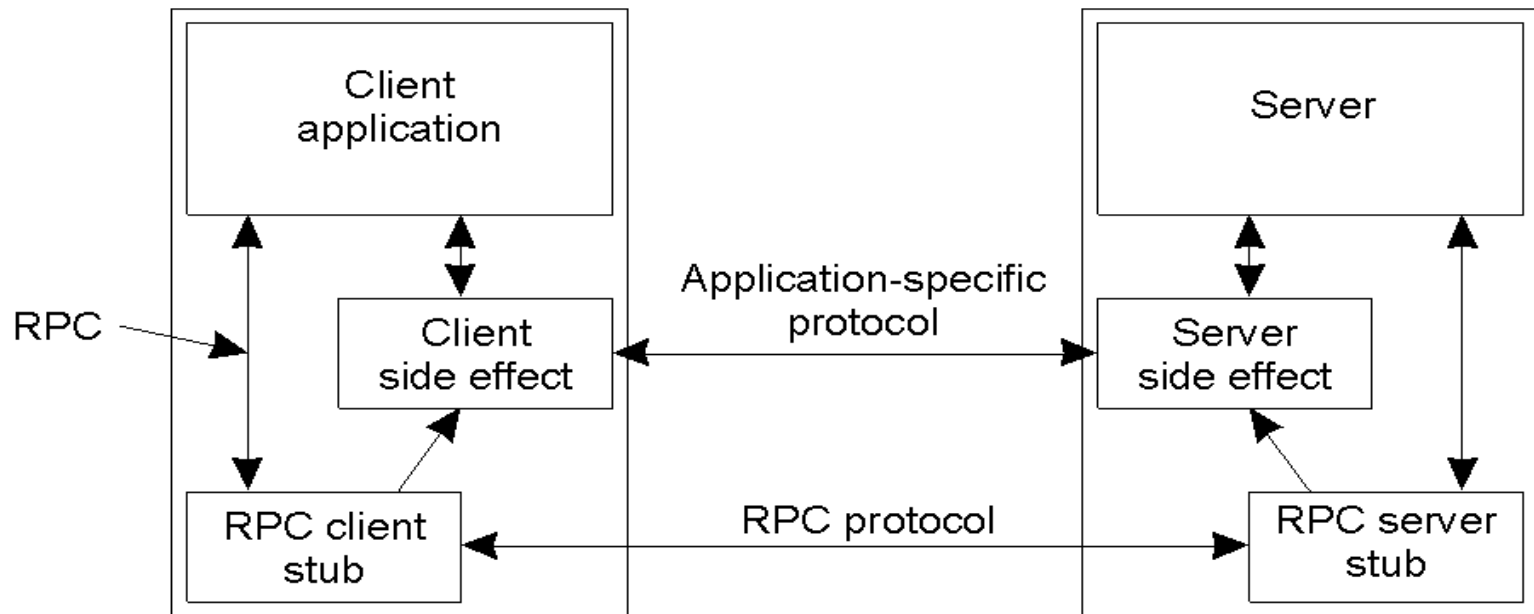
- RPC-based communication (ONC RPC)
  - ↑ RPC hides the differences between distinct OS
- NFS v4 supports **compound procedures**



- Group multiple operations into a single RPC request
  - Handled in order as if sent separately
  - If one fails, the rest terminate
- ↑ Save latency of doing multiple RPCs

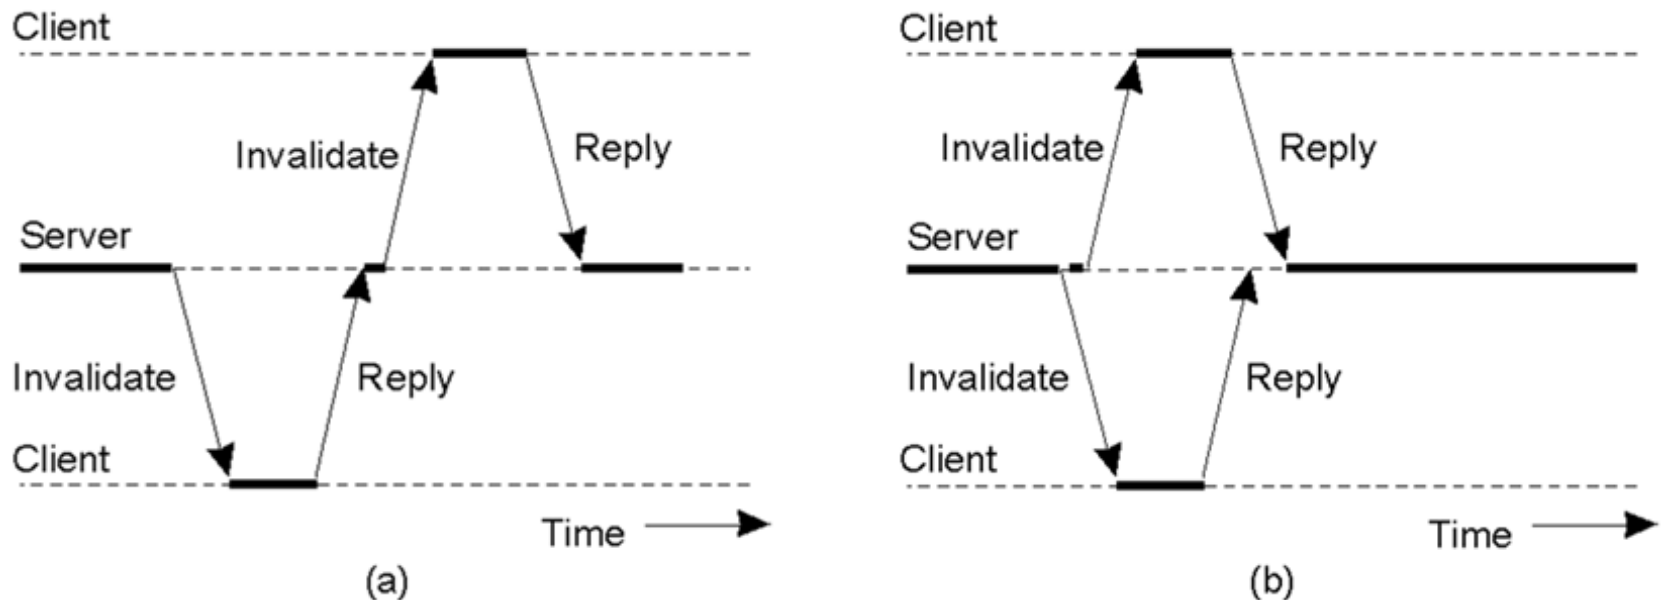
# Coda communication

- Based on **RPC2**
  - Reliable transmission on top of UDP (unreliable)
  - Supports **side-effects**, i.e. user-defined protocols
    - e.g. isochronous mode for a video stream



# Coda communication

- Support for multicast transparent to the client through **MultiRPC** (part of RPC2)
  - Send multiple RPCs in parallel (e.g. invalidations)



Sending invalidation messages one at a time (a) vs. in parallel (b)

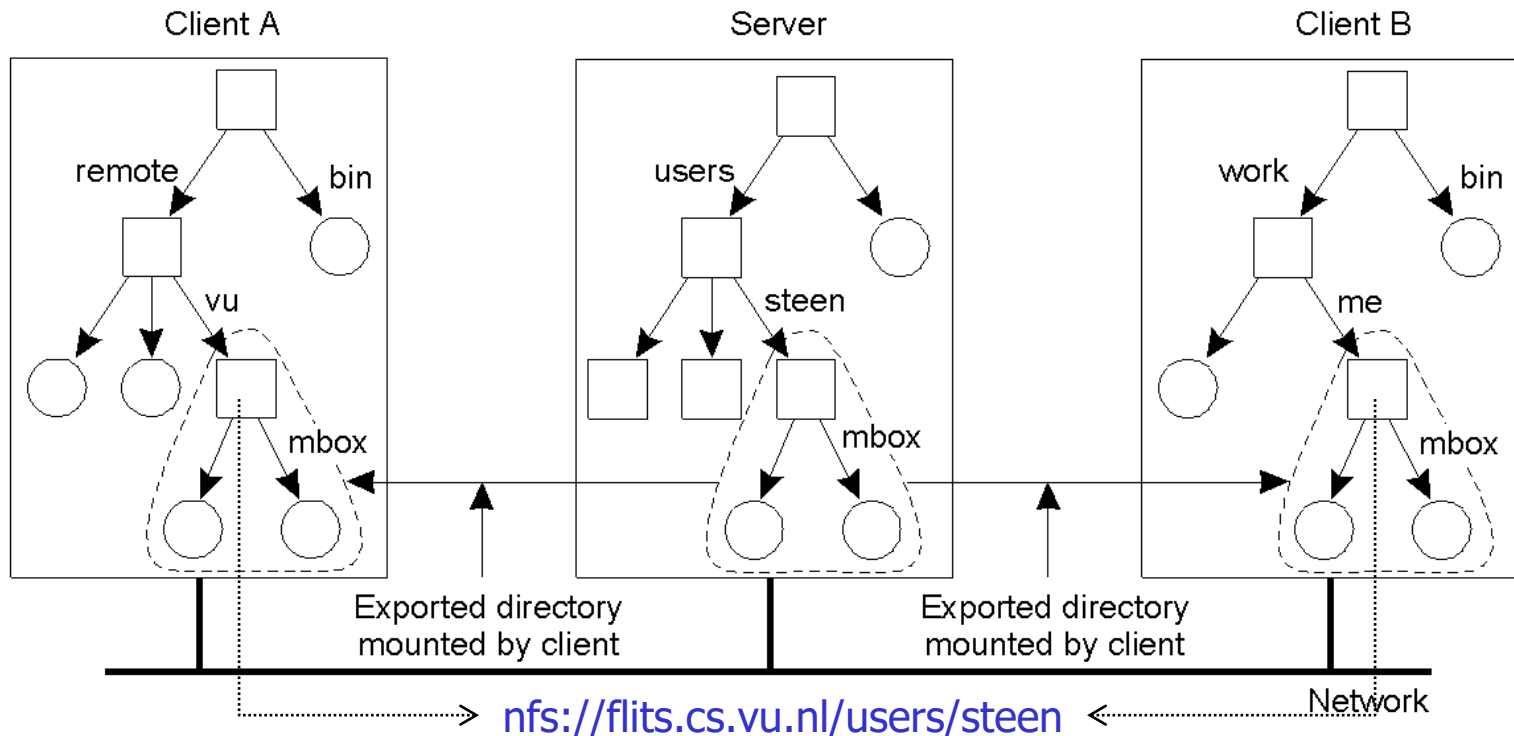
# Contents

---

- Architecture
- Communication
- **Naming**
- Synchronization
- Consistency & replication
- Fault tolerance

# NFS naming

- Transparent access to remote file system
  - By allowing a client to **mount** (part of) a remote file system into its own local file system



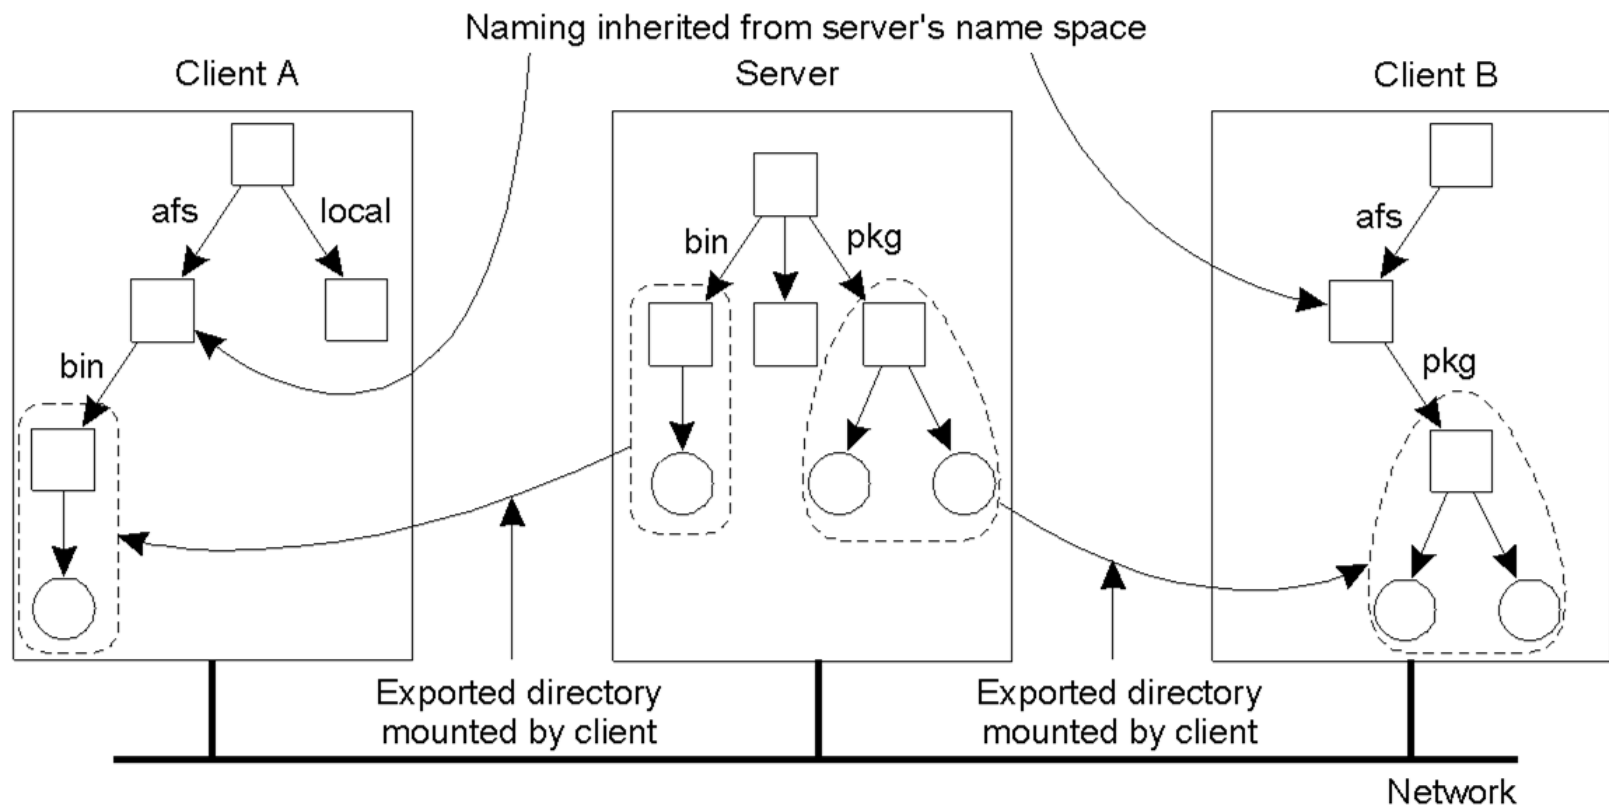
# NFS naming

---

- Pathnames are not globally unique, as they depend on the clients' local name spaces
  - ↓ Referring to shared files is harder
    - ⇒ Partly standardized name space (e.g. /usr/bin)
- Name resolution in NFS v3 is iterative
  - NFS v4 supports recursive lookups within a server
- File handle created by server hosting the file
  - Unique with respect to all its exported file systems
  - File handles must be identifiers
    - File has the same unique file handle during its lifetime, which should not be reused (even after file deletion)

# Coda naming

- Single global shared name space at **/afs**
  - All clients see the same names (unlike NFS)





# Coda naming

---

- File handles are globally unique
  - vs. server unique in NFS
- Files are grouped into **volumes**
  - Each file contained in exactly one volume
  - Volumes may be mounted
    - A volume is the mounting unit
  - Volumes may be replicated among several servers
    - A volume is the unit of server-side replication
  - Distinction between physical and logical volumes
    - A replicated volume consists of several physical volumes that are managed as one logical volume

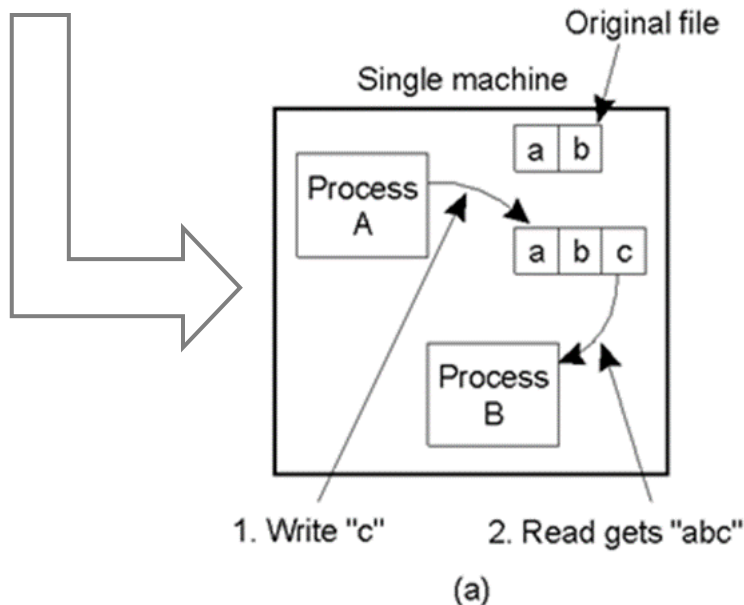
# Contents

---

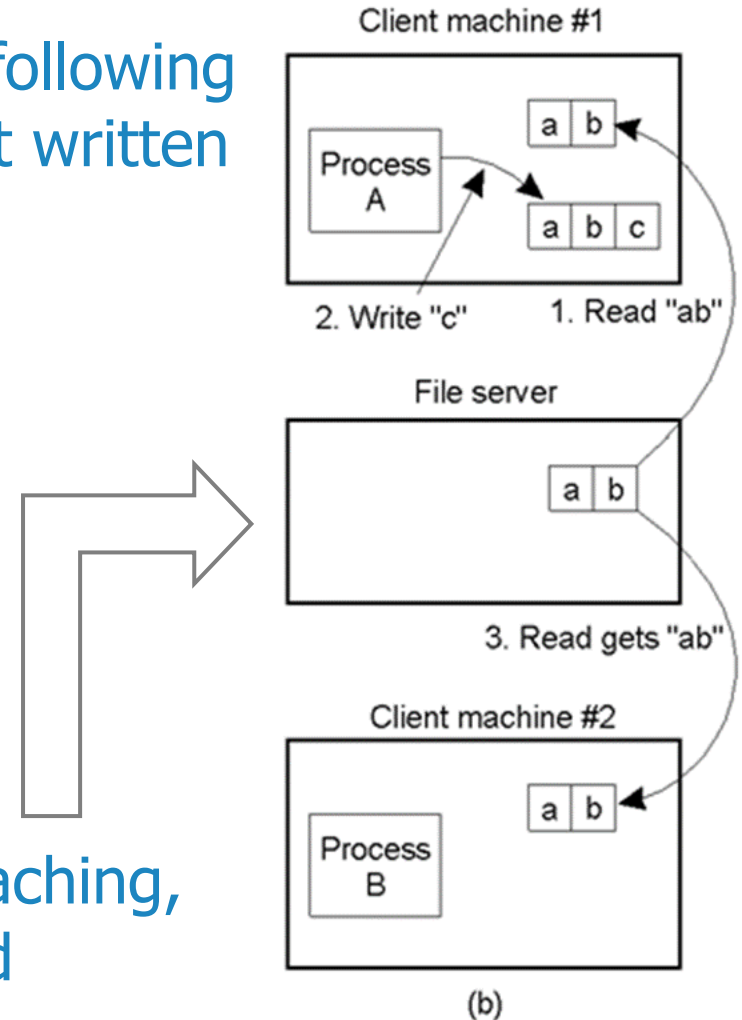
- Architecture
- Communication
- Naming
- **Synchronization**
  - Consistency & replication
  - Fault tolerance

# Semantics of file sharing

- a) In a single machine, a read following a write returns the value just written



- b) In distributed system with caching, stale values may be returned



# Semantics of file sharing

---

- 4 approaches for dealing with shared files

Method	Comment
<b>UNIX semantics</b>	Every operation on a file is instantly visible to all processes
<b>Session semantics</b>	No changes are visible to other processes until the file is closed
<b>Immutable files</b>	No updates are possible, but files can be replaced
<b>Transactions</b>	All changes occur atomically, isolated from other transactions

- NFS can use remote access model for providing UNIX semantics  $\Rightarrow$  performance problem
- Most NFS implementations use local caches for performance implementing **session semantics**
- Coda treats sessions (open...close) as transactions

# File locking

---

- NFS v1-v3: use a separate (stateful) lock manager
- NFS v4: integrated into NFS file access protocol
  - Multiple readers/single writer: read locks vs. write locks
  - Locks granted for a specific time (i.e. lease)

Operation	Description
Lock	Request a lock for a range of bytes (non-blocking)
Lockt	Test whether a conflicting lock has been granted
Locku	Remove a lock from a range of bytes
Renew	Renew the lease on a specified lock

- In addition, **share reservations** can be used
  - Implicit way to lock a file, independent from locks

# File locking

- Result of open operation with share reservations
  - Client requests file access given current denial state
  - Client requests denial state given current file access state

**Current file denial state**

(a)  
**Requested  
file  
access**

	NONE	READ	WRITE	BOTH
READ	Succeed	Fail	Succeed	Fail
WRITE	Succeed	Succeed	Fail	Fail
BOTH	Succeed	Fail	Fail	Fail

**Requested file denial state**

(b)  
**Current  
file  
access  
state**

	NONE	READ	WRITE	BOTH
READ	Succeed	Fail	Succeed	Fail
WRITE	Succeed	Succeed	Fail	Fail
BOTH	Succeed	Fail	Fail	Fail

# Contents

---

- Architecture
- Communication
- Naming
- Synchronization
- **Consistency & replication**
- Fault tolerance

# NFS client-side caching

---

- Preferred technique to attain performance

## a) Data and metadata caching

- Cache the results of read, write, getattr, lookup, and readdir operations
  - i.e. file data, file handles, attributes, directories
- Carry out write operations locally
  - Modified blocks are marked as 'dirty'
  - **Closing** the file (or explicit sync) implies flushing dirty blocks to the server
- Validate cached blocks when they are used using a timestamp-based procedure



# NFS client-side caching

---

- Each cache block has two timestamps
  - $T_c$ : time when cache block was last validated
  - $T_m$ : time when block was last modified at the server
- A block is valid at time  $T$  if either:
  - i.  $T - T_c < F$  (a.k.a. freshness interval)
    - Can be evaluated without accessing the server
    - $F$  is typically 3-30s for files, 30-60s for directories
  - ii.  $T_m \text{ (at server)} = T_m \text{ (at client)}$ 
    - $T_m \text{ (at server)}$  checked only if (i) is false
      - If (ii) is true,  $T_c$  is set to  $T$
      - If (ii) is false, get fresh data from the server and set  $T_m \text{ (at client)}$  and  $T_c$  to  $T$

# NFS client-side caching

---

- Close-to-open cache consistency
  - Extends the previous procedure to support the typical file sharing scenario
    - i.e. sequential file sharing: client A opens a file, writes something and closes it; then client B opens the same file and reads the changes
  - Client forces a cache validity check with the server when the file is opened
    - Ignoring any cache time remaining
  - On closing the file, pending changes are flushed to the server so that the next opener can view them

# NFS client-side caching

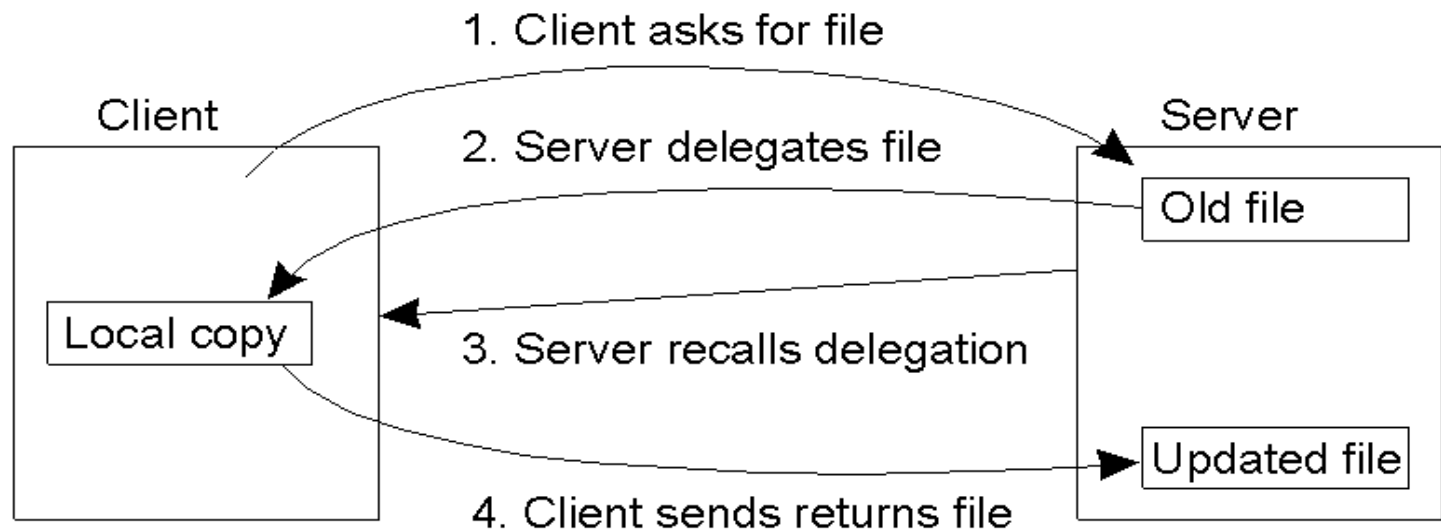
---

## b) NFS v4 supports **open delegation**

- Server voluntarily and temporarily transfers the control of operations on a file to the client
  - Decision is made by the server based on a set of conditions such as the recent history of the file
- During the duration of the delegation, client can treat the file as if no other conflicting accesses are performed by other clients
  - Client can locally service operations such as open, close, or lock without interaction with the server
  - Client can access the file in the cache without sending validation requests to the server

# NFS client-side caching

- Delegations can be recalled by the server using callbacks (i.e. RPC to the client)
  - When another client requests access to the file that conflicts with the granted delegation
  - They have also a lease that is subject to renewal



# NFS client-side caching

---

- Read delegation
  - Awarded on a file opened for reading (not denying read access to others)
    - Multiple read delegations to different clients can be outstanding simultaneously
  - Allows the client to handle locally requests to:
    - Open for reading (not denying read access to others)
    - Read from the cache without checking validity
  - Requests to open the file for writing or to lock the file must still be sent to the server
  - Recalled when another client requests to open the file for writing (or denying read access)

# NFS client-side caching

---

- Write delegation
  - Awarded on a file opened for writing (or r/w)
    - Only one write delegation may exist for a given file at a time, and it is inconsistent with any read delegations
  - While the delegation is outstanding, all the operations in the file can be handled locally
  - Recalled when another client requests to open the file (independently on the requested access)
    - On returning the delegation, the client will commit all dirty data to the server

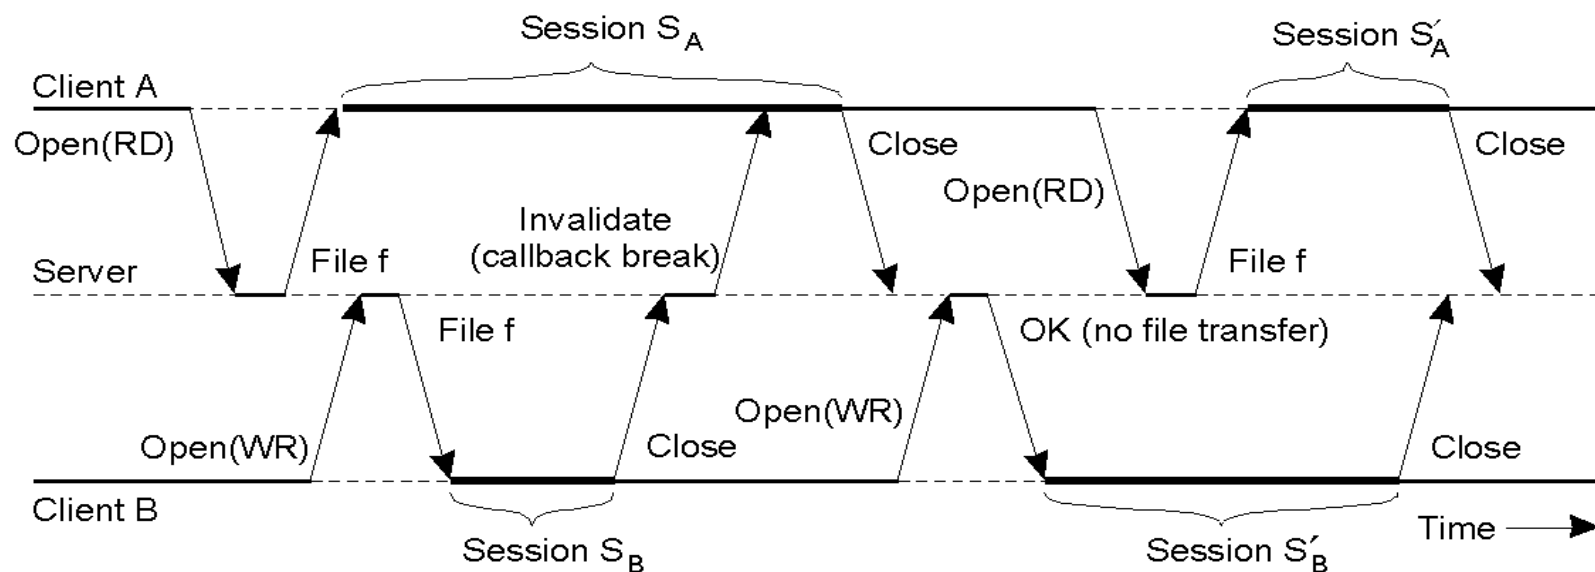
# Coda client-side caching

---

- Coda creates a local copy of the file when it is opened (entire file is cached)
  - All read/write operations done on the local copy
  - Updates are sent to server when the file is closed
- Consistency maintained using **callbacks**
  - On opening a file, client gets a callback promise for the file from the server
  - When another client updates the file, the server breaks the promise (it can also expire by timeout)
    - Server tracks all the clients that have a copy of the file

# Coda client-side caching

- If the promise is valid, a cached file can be re-opened without fetching it from the server
- If the promise has been canceled, re-opening implies fetching a fresh copy of the file






# Coda server-side replication

---

- This is generally used for fault tolerance
- VSG: Volume Storage Group
  - Collection of servers that have a copy of a volume
- AVSG: Accessible VSG
  - Servers in the VSG that the client can contact
  - If AVSG is empty, the client is disconnected
- Replicated-write protocol to keep consistency
  - Variant of ROWA: Read-One Write-All
  - Client opens a file: contact one server in the AVSG
  - Client closes an updated file: multicast file to all the servers in the AVSG using MultiRPC

# Coda server-side replication

- What happens in the presence of partitions?
  - Possible inconsistency: use versioning to detect
    - Version vector when partition happens:  $[1,1,1]$
    - A updates file; version vector in its partition:  $[2,2,1]$
    - B updates file; version vector in its partition:  $[1,1,2]$
    - Partition repaired, compare version vectors: **conflict!**



# Contents

---

- Architecture
- Communication
- Naming
- Synchronization
- Consistency & replication
- **Fault tolerance**

# NFS fault tolerance

---

- Communication failures
  - Most operations are idempotent: can be retried until receiving a reply (at-least-once semantics)
  - Number requests to deal with duplicated non-idempotent operations (as discussed in Lesson 2)
- Client failures (addressed with leases)
  - Server can recover client's locks and delegations when the associated leases have expired
  - Delegations may need to be reclaimed after client restarts to reestablish the file state on the server
    - Client may have unsaved file data stored locally

# NFS fault tolerance


---

- Server failure (addressed with grace periods)
  - When the server restarts, it initiates a grace period of length equal to the lease period
  - During the grace period, clients and server work to reestablish the server state that existed prior to the failure
  - Clients send reclaim-type open and lock requests
    - Corresponding to the locks and open files they had prior to the failure
  - Server rejects new open and lock requests
    - Unless it can reliably determine that granting any such request cannot conflict with a subsequent reclaim

# Coda fault tolerance

---

- Disconnected operation  $\Rightarrow$  high availability
  - HOARDING: Normal operation + cache in advance files that should be available during the disconnection (advised by client)
  - EMULATION: When disconnected, all file requests are serviced using the locally cached copy of the files
  - REINTEGRATION: transfer updates to servers; detect conflicts



# Summary

Issue	NFS	Coda
Design goals	Access transparency	High availability
Access model	Remote	Upload/Download
Communication	RPC (ONC RPC)	RPC (RPC2)
Client process	Thin/Fat	Fat
Server groups	No	Yes
Mount granularity	Directory	File system (volume)
Name space	Per client	Global
File ID scope	File server	Global
Sharing semantics	Session	Transactional
Cache consistency	Write-back	Write-back
Replication	Minimal	ROWA
Fault tolerance	Reliable communication	Replication & caching
Recovery	Client-based	Reintegration

# Summary

---

- Further details:
  - [Tanenbaum]: chapters 2.3.3, 6.3.5, and 7.5.3 (and chapter 11 in 2<sup>nd</sup> Ed.)
  - [Coulouris]: chapters 12, 18.4.3, and 21.5.1